

1 The memoized Held-Karp algorithm

We prepared a memoized Held-Karp algorithm using the top-down dynamic programming approach. This algorithm takes in an undirected, weighted complete graph in the form of an adjacency matrix, as well as a given start vertex. It returns the shortest tour of all vertices in the graph that begin at the start vertex.

1.1 The code

```
1  //////////////////////////////////////
2  ///                               The memoized Held-Karp algorithm          ///
3  //////////////////////////////////////
4
5  // We implement a data structure for storing the visited tours. The first entry
6  // stores a Set of edges in the tour; the second stores the cost of the tour.
7  function StoreTours(edges, cost) {
8      this.edges = edges;
9      this.cost = cost;
10 }
11
12 // For comparing two Sets of edges in (sub-)tours, we implement a check to see
13 // whether they are the same. This was taken from code found at
14 // https://stackoverflow.com/questions/31128855.
15 // Returns TRUE if the sets are the same, FALSE otherwise.
16 function sameSet(set1, set2) {
17     if (set1.size !== set2.size) return false; // obvious and quick check
18     for (let x of set1) if (!set2.has(x)) return false;
19
20     return true;
21 }
22
23
24 // The Held-Karp memoized algorithm, modified from pseudocode in the lectures.
25 // Parameters:
26 // graph => an adjacency matrix for the undirected, weighted graph.
27 // unvisited => a list of unvisited vertices, defaulting to all of them
28 // start => a user-specified vertex from which the tour begins
29 let storedTours = new Array(2); // Store all tours and sub-tours
30     storedTours[0] = new Array();
31     storedTours[1] = new Array();
32
33 function heldKarp(graph, unvisited, start) {
34     // Memoization check
35     for (let i = 0; i < storedTours[0].length; i++) {
36         if (sameSet(unvisited, storedTours[0][i])) {
37             return storedTours[1][i];
38         }
39     }
40
41     if (unvisited.length <= 1) return 0; // No tours to consider
42
43     if (unvisited.length === 2) {
44         let tour = new Set(unvisited);
45         let cost = graph[unvisited[0]][unvisited[1]];
46
47         storedTours[0].push(tour);
48         storedTours[1].push(cost);
49         return cost;
50     }
51
52     else {
53         // Filter out the start vertex
```

```

54   let theRest = unvisited.filter(vert => vert !== start);
55
56   let tour = new Set(theRest);
57   let cost = Infinity;
58   for (let i = 0; i < theRest.length; i++) {
59     let testCost = heldKarp(graph, theRest, theRest[i]) +
60       graph[start][theRest[i]];
61
62     if (testCost < cost) {
63       cost = testCost;
64     }
65   }
66
67   storedTours[0].push(tour);
68   storedTours[1].push(cost);
69   return cost;
70 }
71 }

```

1.2 Empirical time complexity

We investigated the empirical time complexity on Chase's home computer, running a sequence of graphs of increasing size. The results, of number n of vertices against runtime, were plotted in Excel, along with the least-squares trendline plotted. Note that only graphs with 7–10 vertices are plotted, because smaller graphs took less than 1 ms to run (showing up as 0 on our timer) and the 11-vertex graph took more time than we allotted.




Figure 1: The empirical runtime for the Held-Karp algorithm.

Given the empirical runtime trendline shown, we would expect a graph on 11 vertices to take approximately $(2.96 \cdot 10^{-13})e^{(4.19 \times 11)} = 30,755,612$ milliseconds, or about 8.5 hours. This is well above the one-hour threshold.

1.3 Worst-case asymptotic complexity

It is well known that among a set of size n there are 2^n total subsets. This is also the number of tours and sub-tours of a graph with n vertices. In the worst case, the Held-Karp algorithm stores each of them in the cache; furthermore, each one requires storing up to $n + 1$ values (the n vertices visited in the tour and the cost associated with it). So the total space complexity of the Held-Karp implementation is $\Theta(n \times 2^n) = \Theta(2^{n+1})$.

Because of memoization, the algorithm will spend (more than a constant amount of) time on each subset no more than once. Thus, in the worst case, the main body of the algorithm will run 2^n times. The base case of the algorithm (when the size of the subtour is ≤ 2) requires only constant time to run, but the recursive part may require time linear in the number of vertices, so this part of the algorithm requires $\Theta(n \times 2^n) = \Theta(2^{n+1})$ time as well. The memoization check can also take $\Theta(2^n)$ time, since there may be that many entries stored in the cache.

Thus, the overall worst-case time complexity of this implementation of the Held-Karp algorithm is $\Theta(2^{n+1})$.

2 The stochastic 2-opt local search

We also prepared a stochastic 2-opt local search algorithm. This code is not guaranteed to produce an optimal result, but it runs much more quickly than does the Held-Karp implementation. Like the Held-Karp algorithm, the 2-opt stochastic local search algorithm takes an adjacency matrix for an undirected, complete graph and a starting vertex as input and returns an approximate cost for the shortest tour starting at the start vertex.

2.1 The code

```

1  //////////////////////////////////////
2  ///          The 2-opt stochastic local search algorithm          ///
3  //////////////////////////////////////
4
5  // We begin with helper functions: one to randomly reverse part of a given
6  // tour; and one to randomly generate a tour (e.g., to start the algorithm).
7
8
9  // Randomize the elements of an array to find a random starting tour
10 // Found at https://gomakethings.com/how-to-shuffle-an-array-with-vanilla-js/
11 function shuffle(array) {
12     let i = array.length,
13         temp,
14         random_i;
15
16     // While there remain elements to shuffle...
17     while (i !== 0) {
18         //Pick a remaining element...
19         random_i = Math.floor(Math.random() * i);
20         i--;
21
22         //And swap it with the current element.
23         temp = array[i];
24         array[i] = array[random_i];
25         array[random_i] = temp;
26     }
27     return array;
28 }
29
30
31 // Reverse the part of the route between indices i and k, returning the
32 // new route.
33 function two_opt_reversed(route,i,k) {
34
35     // Temporary copy of the input route
36     let copy = route.slice();
37
38     // Reverse the section of the reverse between i and k
39     for (let x = i-1, z=k-1; x < k; x++) {
40         route[z] = copy[x];
41         z--;

```

```
42 }
43 return route;
44 }
45
46 // Stochastic 2-opt local search algorithm running 100*|V|^2 times.
47 // Takes as input an adjacency matrix and a start vertex.
48 function twoOptIter(graph, start) {
49
50     // Small graph corner cases
51     if (graph.length <= 1) return 0;
52
53     // Function global variables
54     let tour = new Array(graph.length),
55         cost = Infinity,
56         maxIter = 100*graph.length*graph.length;
57
58
59     // Generate the random tour (save for the start vertex)
60     for (let i = 0; i < tour.length; i++) tour[i] = i;
61     tour = tour.filter(vert => vert !== start);
62
63     tour = shuffle(tour);
64
65     // Find the initial cost of the tour
66     for (let i = 0; i < tour.length - 1; i++) {
67         cost += graph[tour[i]][tour[i+1]];
68     }
69
70     // Add the distance to the start
71     cost += graph[start][tour[0]];
72
73     // Randomize the route, checking for convergence
74     for (let numIter = 0; numIter < maxIter; numIter++) {
75
76         // Temporary tour cost
77         let tempCost = 0;
78
79         // Boundaries on which to reverse the route
80         let i = Math.ceil(tour.length*Math.random()),
81             k = Math.ceil(tour.length*Math.random());
82
83         // Construct partially reversed tour and find its cost
84         tour = two_opt_reversed(tour, Math.min(i,k), Math.max(i,k));
85
86         for (let i = 0; i < tour.length - 1; i++) {
87             tempCost += graph[tour[i]][tour[i+1]];
88         }
89
90         tempCost += graph[start][tour[0]];
91
92         if (tempCost < cost) cost = tempCost;
93     }
94 }
95
96 return cost;
97 }
```

2.2 Empirical time complexity

We investigated the empirical time complexity on Chase's home computer, running a sequence of graphs of increasing size. The results, of number n of vertices against runtime, were plotted in Excel.

The cubic curve fits the plot very well, and the longest runtime is 4,371,957 ms, or about 1.21 hours.

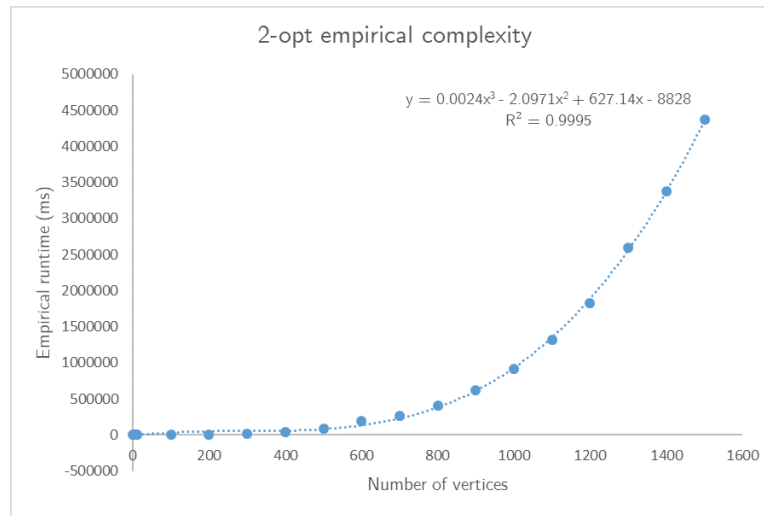


Figure 2: The empirical runtime for the 2-opt graph

2.3 Worst-case asymptotic complexity

First, consider the helper functions. The `shuffle()` function runs in time linear in the input size; namely, the entire initial tour, which has length n for a graph on n vertices. So this function runs in $\Theta(n)$ time, and requires only a constant amount of extra memory.

Likewise, the `two_opt_reversed()` function runs in time linear in the distance between the indices; in the worst case, it may reverse the entire tour array, requiring $\Theta(n)$ time. Additionally, this function creates a temporary array, also requiring $\Theta(n)$ space in the worst case.

Finally, the main function runs exactly $100n^2$ times, where n is the number of vertices in the graph. We chose this number because the number of edges in the graph is on the order of n^2 , so multiplying this by a moderately large constant should give a good chance of a good approximation to the shortest tour. The algorithm calls `shuffle()` exactly once, but calls `two_opt_reversed()` once each iteration; in addition, it requires $\Theta(n)$ time to update the cost each iteration. Apart from the extra space required by its helper functions, the main function requires only a constant amount of extra memory.

As such, the overall runtime of the stochastic 2-opt local search is $\Theta(n + n^2 \times (n + n)) = \Theta(n^3)$, and it requires $\Theta(n)$ space. Running in polynomial time, with a fairly small exponent, and in linear space, this algorithm is much more efficient than Held-Karp's implementation above.

3 Empirical comparison

We plot the empirical runtimes for both algorithms together. Although the 10-vertex graph finishes much more quickly on Held-Karp than the 1500-vertex graph on the 2-opt algorithm, we predict that the 11-vertex graph will take much longer.

We also record the true shortest path (from Held-Karp) and the estimated shortest path (from 2-opt).

Mostly, for these small graphs, the shortest paths found are the same length. However, the 2-opt graph finds slightly longer paths for the graphs of size 9 and 10. This is to be expected, since 2-opt stochastic local search finds only an approximate solution to the shortest tour.

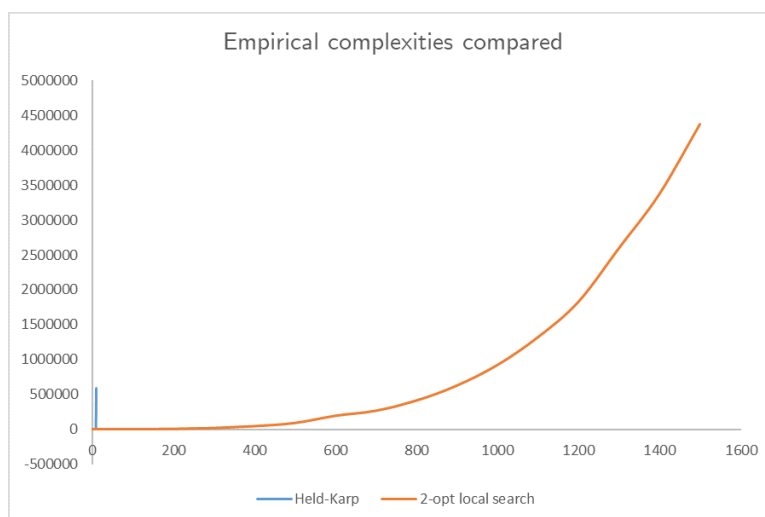


Figure 3: The empirical runtimes plotted together.

Number of vertices	True shortest tour	Approximate shortest tour
0	0	0
1	0	0
2	6	6
3	5	5
4	14	14
5	16	16
6	17	17
7	21	21
8	19	19
9	19	20
10	17	19

Table 1: The shortest paths found.